

Automated-review-rating-system

This project aims to automatically predicts the numerical ratings of the reviews given by the user. It uses Natural Language Processing (NLP), & Deep learning techniques for it.

Project overview and objective


This project aims to build an automated review rating system that analyzes customer reviews and predicts the rating associated with each review. In many real-world scenarios, users provide detailed textual feedback, but numeric ratings may be missing, inconsistent, or biased. The goal is to train a Deep learning model that can understand the sentiment and content of a review and accurately classify it into one of the standard rating categories (e.g., 1 to 5 stars).

By processing and analyzing review text using natural language processing techniques, this system helps businesses quickly identify the sentiment of customer feedback, monitor product performance, and make data-driven decisions without manual intervention.

Environmental Setup

The project was implemented using Python for data processing and analysis due to its rich ecosystem of libraries. Development and testing were carried out in Jupyter Notebook, providing an interactive environment for writing and visualizing code outputs. The entire codebase, along with dataset references and updates, was version-controlled using GitHub for collaboration and backup.

GitHub Project Setup:



Dataset description

To build the automated review rating system, data was gathered in two phases from various sources like Kaggle, Google search datasets, Data web scraping using API Key method in which, containing customer reviews along with corresponding ratings. The objective was to retain only the review text and rating fields after cleaning and preprocessing.

1. Amazon Fashion Reviews (<https://www.kaggle.com/datasets/fawadhossaini1415/amazon-fashion-800k-user-reviews-dataset>)

-This dataset comprises a large-scale collection of Amazon Reviews, gathered in 2023 by the McAuley Lab. Specifically, it focuses on the Amazon Fashion category, containing a total of 800K+ user reviews. It serves as a valuable resource for conducting sentiment analysis.

-Key Fields: Title, Review Text, images, asin, parent_asin, User_id, timestamp, Rating, helpful_vote, target, verified_purchase.

Preprocessing Steps

After collecting and cleaning the datasets, the following initial preprocessing steps were applied to prepare the review text and rating data for model training:

1. Data Loading and Inspection

Each dataset was loaded using pandas and inspected using functions like `.head()`, `.tail()`, `.sample()`, `.shape`, `.columns`, `.info()`, and `.describe()` to understand structure and content.

#Load

```
df = pd.read_csv('data.csv')
```

#View

```
df.head()      # First 5 rows
```

```
df.tail()      # Last 5 rows
```

```
df.sample(3)   # Random 3 rows
```

#Inspect

```
df.shape       # Rows and columns
```

```
df.columns     # Column names
```

```
df.info()      # Data types and nulls
```

```
df.describe()  # Summary stats
```

2. Handling Missing Values

- Columns that had a large number of missing values or were irrelevant to the analysis were dropped.
- Rows with missing entries in the review text or rating columns were removed to maintain data quality.

```
df.isnull().sum()          # Count nulls
df.dropna(inplace=True)    # Drop rows with any nulls
df['x'].fillna(df['x'].mean(), inplace=True) # Fill with mean
```

3. Dropping Irrelevant Columns

- Only the review text and rating columns were retained from each dataset.
- Other metadata like dates, user details, product configuration, etc., were considered unnecessary and discarded.

```
df.drop('x', axis=1, inplace=True) # Drop column
```

4. Duplicate Removal

- All datasets were checked for duplicate rows using `.duplicated()`.
- Any duplicate records found (if any) were removed to avoid data redundancy.

```
df.duplicated().sum() #finding duplicates
df.drop_duplicates(inplace=True) #removing duplicates
```

Review Text Processing

6. Emoji and Special Character Removal

- Emojis, symbols, and special characters were removed using regular expressions to focus only on meaningful textual content.

```
# Remove emojis
text = re.sub(r'[^\w\s,]', '', text.encode('ascii', 'ignore').decode('ascii'))
```

```
# Remove special characters and digits (keep letters and basic punctuation if needed)
text = re.sub(r'[^\A-Za-z\s]', '', text)
```

7. Stop Word Removal

-Common stop words (e.g., “the”, “is”, “and”) were removed using NLTK to reduce noise and retain only informative words.

#defining stopword removal function

```
stop_words = set(stopwords.words('english'))
```

```
punctuations = string.punctuation
```

```
def remove_stopwords(text):
```

```
    # Lowercase
```

```
    text = text.lower()
```

```
    # Remove punctuation
```

```
    text = ''.join([char for char in text if char not in punctuations])
```

```
    # Tokenize and remove stopwords
```

```
    words = text.split()
```

```
    filtered = [word for word in words if word not in stop_words]
```

```
    return ' '.join(filtered)
```

-stopwords used:

-From NLTK

-stop word count= 198

['a', 'about', 'above', 'after', 'again', 'against', 'ain', 'all', 'am', 'an', 'and', 'any', 'are', 'aren', 'aren't', 'as', 'at', 'be', 'because', 'been', 'before', 'being', 'below', 'between', 'both', 'but', 'by', 'can', 'couldn', 'couldn't', 'd', 'did', 'didn', 'didn't', 'do', 'does', 'doesn', 'doesn't', 'doing', 'don', 'don't', 'down', 'during', 'each', 'few', 'for', 'from', 'further', 'had', 'hadn', 'hadn't', 'has', 'hasn', 'hasn't', 'have', 'haven', 'haven't', 'having', 'he', 'he'd', 'he'll', 'he's', 'her', 'here', 'hers', 'herself', 'him', 'himself', 'his', 'how', 'i', 'i'd', 'i'll', 'i'm', 'i've', 'if', 'in', 'into', 'is', 'isn', 'isn't', 'it', 'it'd', 'it'll', 'it's', 'its', 'itself', 'just', 'll', 'm', 'ma', 'me', 'mightn', 'mightn't', 'more', 'most', 'mustn', 'mustn't', 'my', 'myself', 'needn', 'needn't', 'no', 'nor', 'not', 'now', 'o', 'of', 'off', 'on', 'once', 'only', 'or', 'other', 'our', 'ours', 'ourselves', 'out', 'over', 'own', 're', 's', 'same', 'shan', 'shan't', 'she', 'she'd', 'she'll', 'she's', 'should', 'should've', 'shouldn', 'shouldn't', 'so', 'some', 'such', 't', 'than', 'that', 'that'll', 'the', 'their', 'theirs', 'them', 'themselves', 'then', 'there', 'these', 'they', 'they'd', 'they'll', 'they're', 'they've', 'this', 'those', 'through', 'to', 'too', 'under', 'until', 'up', 've', 'very', 'was', 'wasn', 'wasn't', 'we', 'we'd', 'we'll', 'we're', 'we've', 'were', 'weren', 'weren't', 'what', 'when', 'where', 'which', 'while', 'who', 'whom', 'why', 'will', 'with', 'won', 'won't', 'wouldn', 'wouldn't', 'y', 'you', 'you'd', 'you'll', 'you're', 'you've', 'your', 'yours', 'yourself', 'yourselves']

8. Lemmatization

-Words were reduced to their root forms (e.g., “running” → “run”) using WordNetLemmatizer, improving consistency across similar word forms.

#defining lemmatization function

```
lemmatizer = WordNetLemmatizer()
```

```
def lemmatize_text(text):
```

```
    words = text.split()
```

```
    lemmatized = [lemmatizer.lemmatize(word) for word in words]
```

```
    return ' '.join(lemmatized)
```

-Why Lemmatization Over Stemming?

Lemmatization was chosen over stemming because it provides more accurate and meaningful word forms by reducing words to their dictionary base form (lemma), whereas stemming often results in incomplete or non-existent words. For example, "running" becomes "run" through lemmatization, but stemming might reduce it to "runn". This helps retain better context and grammatical correctness, which is especially important for tasks involving sentiment analysis and text understanding.

9. Filtering Out Short and Long Reviews

-Reviews with fewer than 3 words or more than 100 words were removed to retain only meaningful and focused feedback.

Keep only reviews that have word count between 3 and 100

```
filtered_df = df[df['Review'].apply(lambda x: 3 <= len(x.split()) <= 200)].reset_index(drop=True)
```

10. Conflicting Review Removal

-Reviews that did not align with their ratings (e.g., positive text with a negative rating or vice versa) were manually identified and removed.

#checking conflicting reviews

```
from textblob import TextBlob
```

Get sentiment polarity for each review (-1 = negative, 1 = positive)

```
df['Polarity'] = df['Review'].apply(lambda x: TextBlob(x).sentiment.polarity)
```

```

# Optional: classify sentiment direction
def classify_sentiment(p):
    if p > 0.2:
        return 'positive'
    elif p < -0.2:
        return 'negative'
    else:
        return 'neutral'
df['Sentiment'] = df['Polarity'].apply(classify_sentiment)

# Define conflict conditions
conflict_mask = (
    ((df['Rating'] >= 4) & (df['Sentiment'] == 'negative')) |
    ((df['Rating'] <= 2) & (df['Sentiment'] == 'positive'))
)
# Get conflicting reviews for optional inspection
conflicting_reviews = df[conflict_mask]

# Remove conflicting reviews
cleaned_df = df[~conflict_mask].reset_index(drop=True)

```

11. Encode Labels

-convert the rating labels (e.g., 1-star, 2-star, ..., 5-star) from their original categorical/string form into numerical integers so they can be processed by the model.

```

label_encoder = LabelEncoder()
df['label'] = label_encoder.fit_transform(df['Rating'])
num_labels = len(label_encoder.classes_)

```

-LabelEncoder(): Assigns each unique rating a unique integer (e.g., "1 star" → 0, "2 star" → 1, etc.).
-num_labels: Stores the number of unique rating classes for use in the output layer of the LSTM model.

12. Tokenization & Padding

-convert the text reviews into numeric sequences that the LSTM can understand, while ensuring all sequences have the same length.

```

max_vocab_size = 10000
max_seq_len = 128

```

```
tokenizer = Tokenizer(num_words=max_vocab_size, oov_token="<OOV>")
tokenizer.fit_on_texts(train_texts)

train_sequences = tokenizer.texts_to_sequences(train_texts)
test_sequences = tokenizer.texts_to_sequences(test_texts)

train_padded = pad_sequences(train_sequences, maxlen=max_seq_len, padding='post',
truncating='post')

test_padded = pad_sequences(test_sequences, maxlen=max_seq_len, padding='post',
truncating='post')
```

- Tokenization: Maps each word to a unique integer ID.
- OOV Token: <OOV> represents words not seen during training.
- Padding & Truncating: Ensures all reviews are exactly max_seq_len tokens long.

Padding: Adds zeros at the end of shorter sequences.

Truncating: Cuts off longer sequences to fit the maximum length.

13. Create TensorFlow Datasets

- wrap the padded sequences and their labels into a TensorFlow Dataset for efficient batching and shuffling during training.

```
batch_size = 32
```

```
train_dataset = tf.data.Dataset.from_tensor_slices((train_padded,
train_labels)).shuffle(1000).batch(batch_size)
```

```
test_dataset = tf.data.Dataset.from_tensor_slices((test_padded,
test_labels)).batch(batch_size)
```

- Shuffle: Randomizes training examples for better learning.
- batch(batch_size): Groups data into batches for efficient processing by the LSTM model.

Visualizations Used

To better understand the data and guide modeling decisions, several visualizations were created during the exploratory data analysis (EDA) phase:

1. Histogram Plots of Rating Distributions

- Visualized how ratings were distributed across the dataset.
- Helped identify class imbalances (e.g., more 5-star reviews than 1-star).

2. Countplots (Per Rating Class)

- Displayed the number of reviews for each rating class using count plots.
- Useful for comparing rating frequency across different datasets.

3. Pie Charts of Overall Sentiment Distribution

- Ratings were grouped into broader sentiment categories like Positive, Neutral, and Negative.
- Pie charts helped visualize the proportion of each sentiment group in the dataset.

Before and After Cleaning/Balancing Comparisons

- Visualizations were also used to compare the dataset before and after cleaning and balancing. This helped in understanding how much the data was affected by preprocessing steps, especially in terms of class distribution. It clearly showed how the ratings became more evenly distributed after balancing, making the dataset suitable for training.


Final Dataset

-It has 867310 rows and 2 columns of 'Review Text' & 'Rating'.

Done visualization on the final data:

1. Pie Chart

```
rating_counts = df2['Rating'].value_counts().sort_index()
plt.pie(rating_counts, labels=rating_counts.index, autopct='%1.1f%%')
plt.title("Pie Chart")
plt.show()
```



Balanced Dataset

-Balanced dataset created from the final cleaned data, by loading it and again checked for any null values or duplicates for further clarification.

-Balanced it to a target count of 50000 at each rating class depending upon the availability of data in final cleaned one.


```
# Define target size
target_count = 50000

# Downsample each class
balanced_df_final = (
    filtered_df.groupby('Rating')
    .apply(lambda x: x.sample(n=target_count, random_state=42))
    .reset_index(drop=True)
)
```

Done visualization on the Balanced data:

1.Histogram

```
plt.hist(balanced_df_final['Rating'],bins=30)
plt.title('Histogram-Ratings')
plt.show()
```



Sample Reviews Per Rating in Balanced Data:

```
for rating in sorted(balanced_df_final['Rating'].unique()):
    print(f"\n Rating: {int(rating)}")
    samples = balanced_df_final[balanced_df_final['Rating'] == rating].sample(3,
random_state=42)
    for i, review in enumerate(samples['Review'].values, 1):
        print(f"{i}. {review}\n")
```

output:

Rating: 1

1. terra le fit ran ver small wanted return told keep would refund dollar going bother throw away disappointed.
2. looked cute girl picture material thin scratchy feeling wear tank top got paid washed delicate probably wearing time
3. receive dress pictured

Rating: 2

1. love color sea glass pendant way small im disappointed
2. better looking picture thin material baggy certain area expecting probably wear around house
3. product started tearing near bottom first day traveling stitch quality real problem pack

Rating: 3

1. love ring metal soft wish metal durable almost broke said cant wear often would like
2. dress cute received several compliment issue dress came hole pocket got worse would put hand pocket needed outfit spot didnt return definitely get pay husband called disposable dress cheap fabric feel like might tear easily perfect occasion end pocket useless tag im hoping dont ruin attempt hand wash totally cute
3. im size large way tight

Rating: 4

1. work well far worn feel like good compression similar expensive compression sock ive used band top cut calf knee little bit uncomfortable wear day noticeable texture thinner thin sock feel like synthetic fiber opposed softer sock feel like traditional cottonblend sock overall seems like solid compression sock especially price good value overall synthetic fiber feel like durable well see compression hold time overall regret purchase

2. nickel allergy definitely give rash besides cute

3. worn earring number occasion znd received many complement wear clipon earring often hard find pair look good often clipon maker try make look earring pierced earring fine miss unique possibility clipon design affords like many clipons unsure tight clasp would whether ear lobe would hurt little couple hour happy say clasp quite comfortable loosened tightened desired found even clasp almost completely loosened danger might fall first wore party screwed little tightly lobe began hurt little hour soon loosened little problem ive worn better part day without problem cant beat get price

Rating: 5

1. pretty daughter want pair thanks hapoy safe holiday

2. love purse get many compliment definitely worth money

3. super cute fit year old perfectly ordered size toddler

Train-Test-split data for Training

To evaluate the model's performance effectively, the dataset was split into two parts: 80% for training and 20% for testing. The training set was used to fit the model, while the test set helped in assessing its performance on unseen data. This split ensures better generalization and avoids overfitting.

```
from sklearn.model_selection import train_test_split
```

```
train_texts, test_texts, train_labels, test_labels = train_test_split(
    df['Review'].tolist(),
    df['label'].tolist(),
    test_size=0.2,
    stratify=df['label'],
    random_state=42
)
```

-Why Stratified Splitting?

Stratified splitting was used to ensure that the distribution of classes (ratings) remains consistent in both the training and testing sets. This is especially important when dealing with imbalanced datasets, as it helps maintain the proportion of each class across splits, leading to more reliable evaluation and better generalization of the model.

DL Model used- Bidirectional LSTM(Long Short Term Memory)

<https://github.com/sanjuraj-c/automated-review-rating-system/blob/main/Deep%20Learning%20Model%20Used-%20Bidirectional%20LSTM.pdf>

Model Choice and Reason

-For this project, a Bidirectional LSTM (Long Short-Term Memory) network is used because:

-LSTM networks are excellent at handling sequential data such as text, as they can remember long-term dependencies in the sequence.

-Bidirectional LSTM processes the input text in both forward and backward directions, allowing the model to capture context from both past and future words in a review. This is especially important for understanding complex sentence structures and nuanced meanings in customer feedback.

-Compared to a standard LSTM, the bidirectional version provides richer context understanding which improves classification accuracy for multi-class rating predictions.

Layer Details:

1.Embedding Layer

-Converts word IDs into dense vectors of size 128.

-Learns meaningful word representations during training.

2.First Bidirectional LSTM (64 units)

-Processes the text in both directions.

-return_sequences=True ensures the entire sequence output is passed to the next LSTM layer.

3.Second Bidirectional LSTM (32 units)

-Further processes the combined sequence data into a fixed-size vector.

4.Dense Layer (64 units, ReLU activation)

-Learns complex feature interactions from LSTM outputs.

5.Dropout Layer (0.5)

-Randomly drops 50% of neurons during training to reduce overfitting.

6.Output Layer

-Dense(num_labels, activation='softmax') outputs a probability distribution over the rating classes.

Model Training Logic

1. Prepare Datasets

-Load and preprocess training and validation datasets.

-Ensure input sequences are tokenized, padded, and matched to the model's expected input shape.

2. Define Early Stopping

-Set up EarlyStopping to monitor validation loss (val_loss).

-Stop training if the loss does not improve for 3 consecutive epochs.

-Restore the best weights from the most accurate epoch.

3. Compile the Model

-Use Categorical Crossentropy as the loss function (multi-class classification).

-Use Adam optimizer for adaptive learning rate optimization.

-Track accuracy as the main evaluation metric.

4. Handle Class Imbalance

-Compute class weights based on the distribution of ratings.

-Pass class_weight to model.fit() to give more importance to underrepresented classes.

5. Train the Model

- Train on the training dataset with a batch size of 32.
- Validate on the test dataset after each epoch.
- Run for up to 20 epochs, or until early stopping triggers.

6. Store Training History

- Save the history object from `model.fit()` to later visualize training and validation accuracy/loss trends.

7. Select Best Model

- Due to early stopping, the final model uses the weights from the epoch where validation loss was lowest.

Evaluation results and interpretation

1.Accuracy: 52.74%

2. Precision/Recall/F1-score:

Classification Report:					
	precision	recall	f1-score	support	
1.0	0.60	0.60	0.60	10000	
2.0	0.42	0.51	0.46	10000	
3.0	0.44	0.35	0.39	10000	
4.0	0.52	0.42	0.47	10000	
5.0	0.65	0.76	0.70	10000	
accuracy			0.53	50000	
macro avg	0.52	0.53	0.52	50000	
weighted avg	0.52	0.53	0.52	50000	

1.Precision: Ranges from 0.42 (lowest for rating 2) to 0.65 (highest for rating 5).

-High precision for rating 5 indicates that when the model predicts a rating of 5, it is often correct.

-Low precision for rating 2 suggests many false positives for this class.


2.Recall: Ranges from 0.35 (lowest for rating 3) to 0.76 (highest for rating 5).

-High recall for rating 5 means the model is able to identify most of the actual 5-star reviews.

-Low recall for rating 3 indicates many 3-star reviews are misclassified into other ratings.

3.F1-Score: Balanced measure of precision and recall; highest for rating 5 (0.70) and lowest for rating 3 (0.39).

3.Confusion matrix



-Rating 5: Most accurately predicted; majority classified correctly (7,641 correct out of 10,000).

-Rating 1: Moderate performance; 5,960 correctly classified but a large portion confused with rating 2.

-Rating 3: Most challenging class; predictions are spread across ratings 2, 4, and 5.

-Rating 2 & 4: Often confused with each other, indicating possible similarity in textual patterns.

